



# 10-601 Introduction to Machine Learning

Machine Learning Department  
School of Computer Science  
Carnegie Mellon University

# Optimization for ML

+

# Logistic Regression

Matt Gormley  
Lecture 8  
Feb. 11, 2019



# 10-601 Introduction to Machine Learning

Machine Learning Department  
School of Computer Science  
Carnegie Mellon University

# Optimization for ML

+

# Logistic Regression

# Probabilistic Learning

Matt Gormley  
Lecture 8  
Feb. 11, 2019

# Q&A

# Reminders

- **Homework 3: KNN, Perceptron, Lin.Reg.**
  - Out: Wed, Feb 6
  - Due: Fri, Feb 15 at 11:59pm
- **Today's In-Class Poll**
  - <http://p8.mlcourse.org>


# **CONVEXITY**

# Convexity

Function  $f : \mathbb{R}^M \rightarrow \mathbb{R}$  is **convex**

if  $\forall \mathbf{x}_1 \in \mathbb{R}^M, \mathbf{x}_2 \in \mathbb{R}^M, 0 \leq t \leq 1$ :

$$f(t\mathbf{x}_1 + (1 - t)\mathbf{x}_2) \leq tf(\mathbf{x}_1) + (1 - t)f(\mathbf{x}_2)$$




# Convexity

Suppose we have a function  $f(x) : \mathcal{X} \rightarrow \mathcal{Y}$ .


- The value  $x^*$  is a **global minimum** of  $f$  iff  $f(x^*) \leq f(x), \forall x \in \mathcal{X}$ .
- The value  $x^*$  is a **local minimum** of  $f$  iff  $\exists \epsilon$  s.t.  $f(x^*) \leq f(x), \forall x \in [x^* - \epsilon, x^* + \epsilon]$ .

## Convex Function



- Each **local minimum** is a **global minimum**

## Nonconvex Function




- A nonconvex function is **not convex**
- Each **local minimum** is **not necessarily a global minimum**

# Convexity

Function  $f : \mathbb{R}^M \rightarrow \mathbb{R}$  is **convex**  
if  $\forall \mathbf{x}_1 \in \mathbb{R}^M, \mathbf{x}_2 \in \mathbb{R}^M, 0 \leq t \leq 1$ :


$$f(t\mathbf{x}_1 + (1-t)\mathbf{x}_2) \leq tf(\mathbf{x}_1) + (1-t)f(\mathbf{x}_2)$$



Each **local minimum** of a **convex** function is also a **global minimum**.

Function  $f : \mathbb{R}^M \rightarrow \mathbb{R}$  is **strictly convex**  
if  $\forall \mathbf{x}_1 \in \mathbb{R}^M, \mathbf{x}_2 \in \mathbb{R}^M, 0 \leq t \leq 1$ :

$$f(t\mathbf{x}_1 + (1-t)\mathbf{x}_2) < tf(\mathbf{x}_1) + (1-t)f(\mathbf{x}_2)$$



A **strictly convex** function has a **unique global minimum**.

# Convexity

The **Mean Squared Error** function,  
which we minimize for learning  
the parameters of Linear  
Regression, is **convex!**

# Regression Loss Functions

## In-Class Exercise:

Which of the following could be used as loss functions for training a linear regression model?

Select all that apply.

- A.  $\ell(\hat{y}, y) = \|\hat{y} - y\|_2$
- B.  $\ell(\hat{y}, y) = |\hat{y} - y|$
- C.  $\ell(\hat{y}, y) = \frac{1}{2}(\hat{y} - y)^2$
- D.  $\ell(\hat{y}, y) = \frac{1}{4}(\hat{y} - y)^4$
- E.  $\ell(\hat{y}, y) = \begin{cases} \frac{1}{2}(\hat{y} - y)^2 & \text{if } |\hat{y} - y| \leq \delta \\ \delta|\hat{y} - y| - \frac{1}{2}\delta^2 & \text{otherwise} \end{cases}$
- F.  $\ell(\hat{y}, y) = \log(\cosh(\hat{y} - y))$

# Solving Linear Regression

## Question:

**True or False:** If Mean Squared Error (i.e.  $\frac{1}{N} \sum_{i=1}^N (y^{(i)} - h(\mathbf{x}^{(i)}))^2$ ) has a unique minimizer (i.e. argmin), then Mean Absolute Error (i.e.  $\frac{1}{N} \sum_{i=1}^N |y^{(i)} - h(\mathbf{x}^{(i)})|$ ) must also have a unique minimizer.

## Answer:

The Big Picture

# **OPTIMIZATION FOR ML**



# Function Approximation

*Chalkboard*

– The Big Picture

# **GRADIENT DESCENT**

# Motivation: Gradient Descent

To solve the Ordinary Least Squares problem we compute:

$$\begin{aligned}\hat{\theta} &= \underset{\theta}{\operatorname{argmin}} = \frac{1}{N} \sum_{i=1}^N \frac{1}{2}(y^{(i)} - (\theta^T \mathbf{x}^{(i)}))^2 \\ &= (\mathbf{X}^T \mathbf{X})^{-1} (\mathbf{X}^T \mathbf{Y})\end{aligned}$$

The resulting shape of the matrices:

$$(\underbrace{\mathbf{X}^T}_{M \times N} \underbrace{\mathbf{X}}_{N \times M})^{-1} (\underbrace{\mathbf{X}^T}_{M \times N} \underbrace{\mathbf{Y}}_{N \times 1})$$

$M \times M$        $M \times 1$

**Background: Matrix Multiplication** Given matrices **A** and **B**

- If **A** is  $q \times r$  and **B** is  $r \times s$ , computing **AB** takes  $O(qrs)$
- If **A** and **B** are  $q \times q$ , computing **AB** takes  $O(q^{2.373})$
- If **A** is  $q \times q$ , computing  $A^{-1}$  takes  $O(q^{2.373})$ .

Computational Complexity of OLS:

$\mathbf{X}^T \mathbf{X}$	$O(M^2 N)$
$(\quad)^{-1}$	$O(M^{2.373})$
$\mathbf{X}^T \mathbf{Y}$	$O(MN)$
$(\quad)^{-1} (\quad)$	$O(M^2)$
total	$O(M^2 N + M^{2.373})$

Linear in # of examples, N  
Polynomial in # of features, M



# Motivation: Gradient Descent

Cases to consider gradient descent:


1. What if we **can not** find a closed-form solution?
2. What if we **can**, but it's inefficient to compute?
3. What if we **can**, but it's numerically unstable to compute?

# Topographical Maps


# Topographical Maps



# Gradients




# Gradients




These are the **gradients** that  
Gradient **Ascent** would follow.

# (Negative) Gradients



These are the **negative** gradients that  
Gradient **Descent** would follow.


# (Negative) Gradient Paths



Shown are the **paths** that Gradient Descent would follow if it were making **infinitesimally small steps**.

# Pros and cons of gradient descent

- Simple and often quite effective on ML tasks
- Often very scalable
- Only applies to smooth functions (differentiable)
- Might find a local minimum, rather than a global one



# Gradient Descent


## *Chalkboard*

- Gradient Descent Algorithm
- Details: starting point, stopping criterion, line search

# Gradient Descent

## Algorithm 1 Gradient Descent

```
1: procedure GD( $\mathcal{D}$ ,  $\theta^{(0)}$ )
2:    $\theta \leftarrow \theta^{(0)}$ 
3:   while not converged do
4:      $\theta \leftarrow \theta - \lambda \nabla_{\theta} J(\theta)$ 
5:   return  $\theta$ 
```




In order to apply GD to Linear Regression all we need is the **gradient** of the objective function (i.e. vector of partial derivatives).

$$\nabla_{\theta} J(\theta) = \begin{bmatrix} \frac{d}{d\theta_1} J(\theta) \\ \frac{d}{d\theta_2} J(\theta) \\ \vdots \\ \frac{d}{d\theta_M} J(\theta) \end{bmatrix}$$

# Gradient Descent

## Algorithm 1 Gradient Descent

```
1: procedure GD( $\mathcal{D}$ ,  $\theta^{(0)}$ )
2:    $\theta \leftarrow \theta^{(0)}$ 
3:   while not converged do
4:      $\theta \leftarrow \theta - \lambda \nabla_{\theta} J(\theta)$ 
5:   return  $\theta$ 
```



There are many possible ways to detect **convergence**.  
For example, we could check whether the L2 norm of  
the gradient is below some small tolerance.

$$\|\nabla_{\theta} J(\theta)\|_2 \leq \epsilon$$

Alternatively we could check that the reduction in the  
objective function from one iteration to the next is small.

# **GRADIENT DESCENT FOR LINEAR REGRESSION**

# Optimization for Linear Regression

## Chalkboard


- Computing the gradient for Linear Regression
- Gradient Descent for Linear Regression
- 2D Example in Three Parts:
  1. Line over time
  2. Parameters space over time
  3. Train / test error over time

# **STOCHASTIC GRADIENT DESCENT**

# Gradient Descent

## Algorithm 1 Gradient Descent


```
1: procedure GD( $\mathcal{D}$ ,  $\theta^{(0)}$ )
2:    $\theta \leftarrow \theta^{(0)}$ 
3:   while not converged do
4:      $\theta \leftarrow \theta - \lambda \nabla_{\theta} J(\theta)$ 
5:   return  $\theta$ 
```



# Stochastic Gradient Descent (SGD)

## Algorithm 2 Stochastic Gradient Descent (SGD)

```
1: procedure SGD( $\mathcal{D}$ ,  $\theta^{(0)}$ )
2:    $\theta \leftarrow \theta^{(0)}$ 
3:   while not converged do
4:     for  $i \sim \text{Uniform}(\{1, 2, \dots, N\})$  do
5:        $\theta \leftarrow \theta - \lambda \nabla_{\theta} J^{(i)}(\theta)$ 
6:   return  $\theta$ 
```




We need a per-example objective:

$$\text{Let } J(\theta) = \sum_{i=1}^N J^{(i)}(\theta)$$

# Stochastic Gradient Descent (SGD)

## Algorithm 2 Stochastic Gradient Descent (SGD)

```
1: procedure SGD( $\mathcal{D}$ ,  $\theta^{(0)}$ )
2:    $\theta \leftarrow \theta^{(0)}$ 
3:   while not converged do
4:     for  $i \sim \text{Uniform}(\{1, 2, \dots, N\})$  do
5:        $\theta \leftarrow \theta - \lambda \nabla_{\theta} J^{(i)}(\theta)$ 
6:   return  $\theta$ 
```




We need a per-example objective:

$$\text{Let } J(\theta) = \sum_{i=1}^N J^{(i)}(\theta)$$

# Stochastic Gradient Descent (SGD)

## Algorithm 2 Stochastic Gradient Descent (SGD)

```
1: procedure SGD( $\mathcal{D}$ ,  $\theta^{(0)}$ )
2:    $\theta \leftarrow \theta^{(0)}$ 
3:   while not converged do
4:     for  $i \in \text{shuffle}(\{1, 2, \dots, N\})$  do
5:        $\theta \leftarrow \theta - \lambda \nabla_{\theta} J^{(i)}(\theta)$ 
6:   return  $\theta$ 
```




We need a per-example objective:


$$\text{Let } J(\theta) = \sum_{i=1}^N J^{(i)}(\theta)$$

In practice, it is common to implement SGD using sampling **without** replacement (i.e.  $\text{shuffle}(\{1, 2, \dots, N\})$ ), even though most of the theory is for sampling **with** replacement (i.e.  $\text{Uniform}(\{1, 2, \dots, N\})$ ).

# Convergence Curves



- Def: an **epoch** is a single pass through the training data
  1. For GD, only **one update** per epoch
  2. For SGD,  **$N$  updates** per epoch  
 $N = (\# \text{ train examples})$
- SGD reduces MSE much more rapidly than GD
- For GD / SGD, training MSE is initially large due to uninformed initialization



# Expectations of Gradients

$$\frac{\partial \bar{J}(\vec{\theta})}{\partial \theta_j} = \frac{1}{N} \sum_{i=1}^N \frac{d}{d\theta_j} (J_i(\vec{\theta}))$$

$$\nabla \bar{J}(\vec{\theta}) = \begin{bmatrix} \vdots \\ j^{\text{th}} \\ \vdots \\ \cdot \end{bmatrix} = \frac{1}{N} \sum_{i=1}^N \nabla J_i(\vec{\theta})$$

Recall : For any discrete r.v.  $X$

$$E_X[f(x)] \triangleq \sum_x P(X=x) f(x)$$

Q: What is the expected value of a randomly chosen  $\nabla J_i(\vec{\theta})$ ?

$$\text{Let } I \sim \text{Uniform}\{\{1, \dots, N\}\}$$

$$\Rightarrow P(I=i) = \frac{1}{N} \text{ if } i \in \{1, \dots, N\}$$

$$\begin{aligned} E_I[\nabla J_I(\vec{\theta})] &= \sum_{i=1}^N P(I=i) \nabla J_i(\vec{\theta}) \\ &\Rightarrow \frac{1}{N} \sum_{i=1}^N \nabla J_i(\vec{\theta}) \\ &= \nabla \bar{J}(\vec{\theta}) \end{aligned}$$

# Convergence of Optimizers

Convergence Analysis:

Def: convergence is when  $J(\vec{\theta}) - J(\vec{\theta}^*) < \epsilon$

Methods	Steps to Converge	Computation per iteration
Newton's Method	$O(\ln \ln \frac{1}{\epsilon})$	$\nabla J(\theta) \quad \nabla^2 J(\theta) \leftarrow O(NM^2)$
GD	$O(\ln \frac{1}{\epsilon})$	$\nabla J(\theta) \leftarrow O(NM)$
SGD	$O(\frac{1}{\epsilon})$	$\nabla J_i(\theta) \leftarrow O(M)$

"almost sure" convergence  
lots of caveats  
and conditions

true unknown min

way less computation

Takeaway: SGD has much slower asymptotic convergence.  
but is often faster in practice.

# Optimization Objectives

*You should be able to...*

- Apply gradient descent to optimize a function
- Apply stochastic gradient descent (SGD) to optimize a function
- Apply knowledge of zero derivatives to identify a closed-form solution (if one exists) to an optimization problem
- Distinguish between convex, concave, and nonconvex functions
- Obtain the gradient (and Hessian) of a (twice) differentiable function

# Linear Regression Objectives

You should be able to...

- Design k-NN Regression and Decision Tree Regression
- Implement learning for Linear Regression using three optimization techniques: (1) closed form, (2) gradient descent, (3) stochastic gradient descent
- Choose a Linear Regression optimization technique that is appropriate for a particular dataset by analyzing the tradeoff of computational complexity vs. convergence speed
- Distinguish the three sources of error identified by the bias-variance decomposition: bias, variance, and irreducible error.